Find the Promos

Kuku AI Generalist assignment | Problem 1 | Solution

Evidence files: [B] Kuku assignment brief; [P] STAGE2_REAL_RUN.md; [R] run_A/eval_report.md and run_state.json; [T] run_A/responses/*.success.body; [A] run_A_user_audit.csv; [S] stability_report.md and run_B/run_state.json; [F] WHERE_AI_GOT_IT_WRONG.md and stage3_results(1).json.

Why this problem and setup

I chose this problem because it addresses Kuku's promo-selection workflow, supplies real input, and produces timestamps I can check [B]. I preferred that scope to a legal assistant requiring specialist grounding or an open-ended problem requiring separate validation. I chose local Ollama for a simple setup and control of the model configuration. The recorded runs use Qwen3:4b-instruct; the model is swappable through configuration [R].

How my tool works

The code parses timed screenplay events and creates a limited menu of candidate time ranges. It checks for 20-60 second duration, ordering, cut boundaries and a source-backed opening in the first 3 seconds [B, P]. Qwen selects and explains promising ranges. The same model judges them in a separate pass. Code combines those judgments with rule-based signals, removes overlapping cuts and adjusts for diversity before ranking. The workflow also tries prevalidated cross-episode combinations.

I use Streamlit to load input, run the pipeline, inspect ranked results and record my review. A structural pass checks timing and format; it does not prove a strong hook or cliffhanger.

What the actual run produced

From [P], I parsed 50 episodes, 100 scenes and 3,727 timed events. From [R], Run A processed 50/50 episodes: 115 proposals, 100 structurally valid, 85 judged, 15 unjudged and 81 ranked after overlap removal. The run is partial.

My manual review

From [A], I reviewed the top 10: Good: ranks 2, 4, 6, 10; Unsure: 3, 7, 8; Bad: 1, 5, 9. Some Good choices still need stronger openings or context. I found ranks 3 and 9 similar in scene and theme. This is one small, non-random script-based review, not evidence of audience performance.

Where AI and my checks fell short

From [F, P], an early ChatGPT-assisted scan counted 98 scenes because it missed uppercase scene labels in Episode 6; a case-insensitive scan found 100. A local-model judge also treated a reply after a cut as if it were inside the clip. I revised the source excerpts in the judging input to show only the selected material, but that does not guarantee factual accuracy.

From [A, F], I found that rank 5's named cliffhanger actually occurs near its beginning, while the clip ends on a resolved response. For rank 8, the explanation describes an opening strike that ends before the clip begins. Rank 3 ends between timed blocks in a continuing thought: a safe event boundary can still be a poor editorial cut.

Run-to-run stability

From [S], the two partial top-10 lists share 3 exact cuts: 30% of each list. Jaccard is 3 shared cuts out of 17 distinct cuts across both lists, or 17.6%. Exact matching misses similar alternative cuts. Both runs used the same local model at temperature 0.7 [R, S], but their labels also changed the prompts, so these results do not establish why the cuts differ. I treat the output as an editor's candidate pool; final selection needs human judgment.

Scale and cost

From [R, T], the table uses Run A's whole-pipeline usage, including cross-episode and judging calls. I divide by the 50 episodes and double the totals for a simple 100-episode projection.

Measure	Run A: 50 episodes	Per episode	Estimate: 100 episodes
Input tokens	534,047	10,681	1,068,094
Output tokens	58,606	1,172	117,212
Model calls	107	2.14	214
Model processing	87.8 min	1.76 min	175.6 min

The processing total is the sum of provider-reported call durations, not end-to-end wall time. The estimate assumes similar inputs, hardware and outcomes; Run A was partial. Local inference has no per-token charge, but hardware and electricity still cost money.

Proposed, not implemented: queue shows, process episodes across suitable workers, then combine and rank candidates for editorial review. New script formats need validation.

Limits and AI use

The same model selects and judges; that is not independent validation. I cannot judge acting, music, visual impact or final-video timing from scripts. With more time, I would inspect actual footage and test a stronger or independent model; neither is claimed here. I used ChatGPT for exploration, code, prompts and debugging [F].